No training: with KV crumbs of the writer’s own sentence the answer picks right 77% (0.6B) and 87% (4B) of the time, against 14% and 17% with crumbs cut

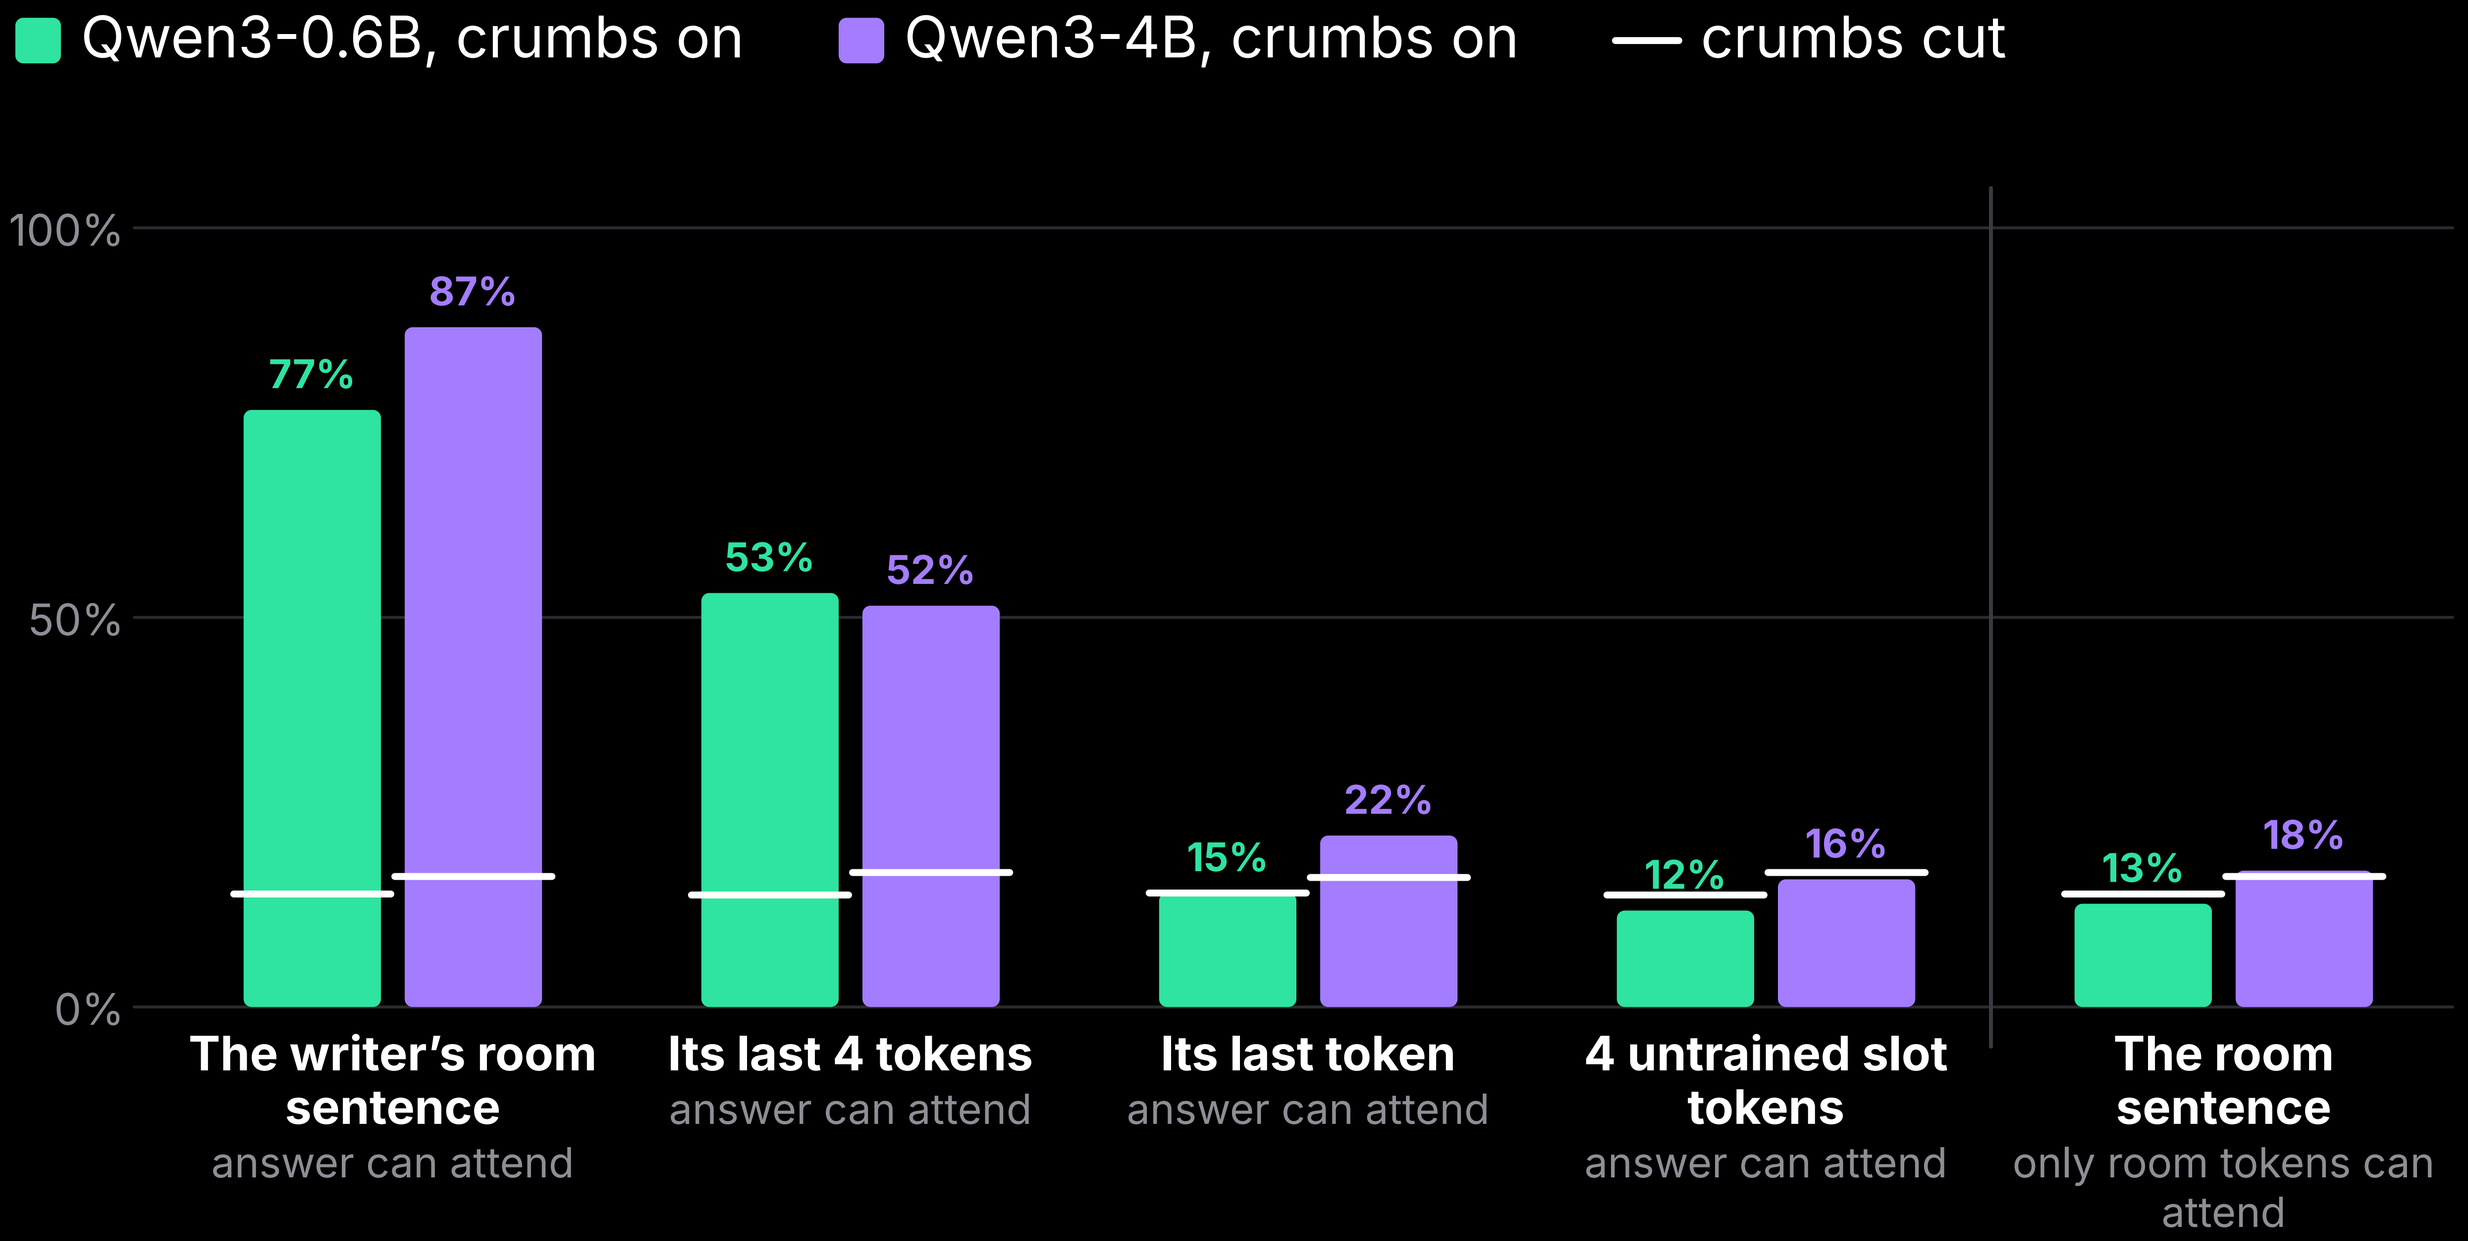

Untrained models, no adapter: the crumbs are the writer’s own keys and values in every layer at its own positions. Answer can attend: every follower token, the answer’s included, attends to the crumbs of each room visited so far; otherwise only the follower’s tokens in a room attend to that room’s crumbs. Forced choice on trail: each visited room’s object against the episode’s 5 other objects and the key’s room against the 5 other rooms, 200 test episodes per bar.