A penalty on letters broke the messages between two models before any code formed

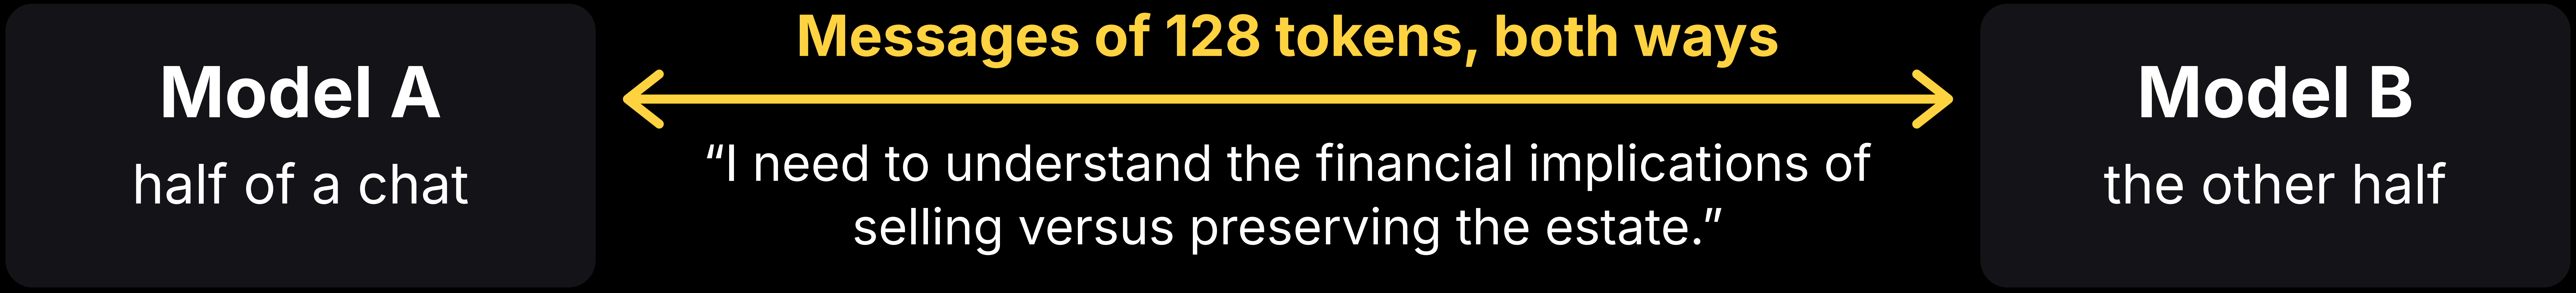

Model A
half of a chat

Messages of 128 tokens, both ways

"I need to understand the financial implications of selling versus preserving the estate."

Model B
the other half

Goal: each model answers like a model that read the whole chat

Each row is one real message, each cell one token: letters, other symbols, blank ▭

Weak penalty on letters, 1,400 steps

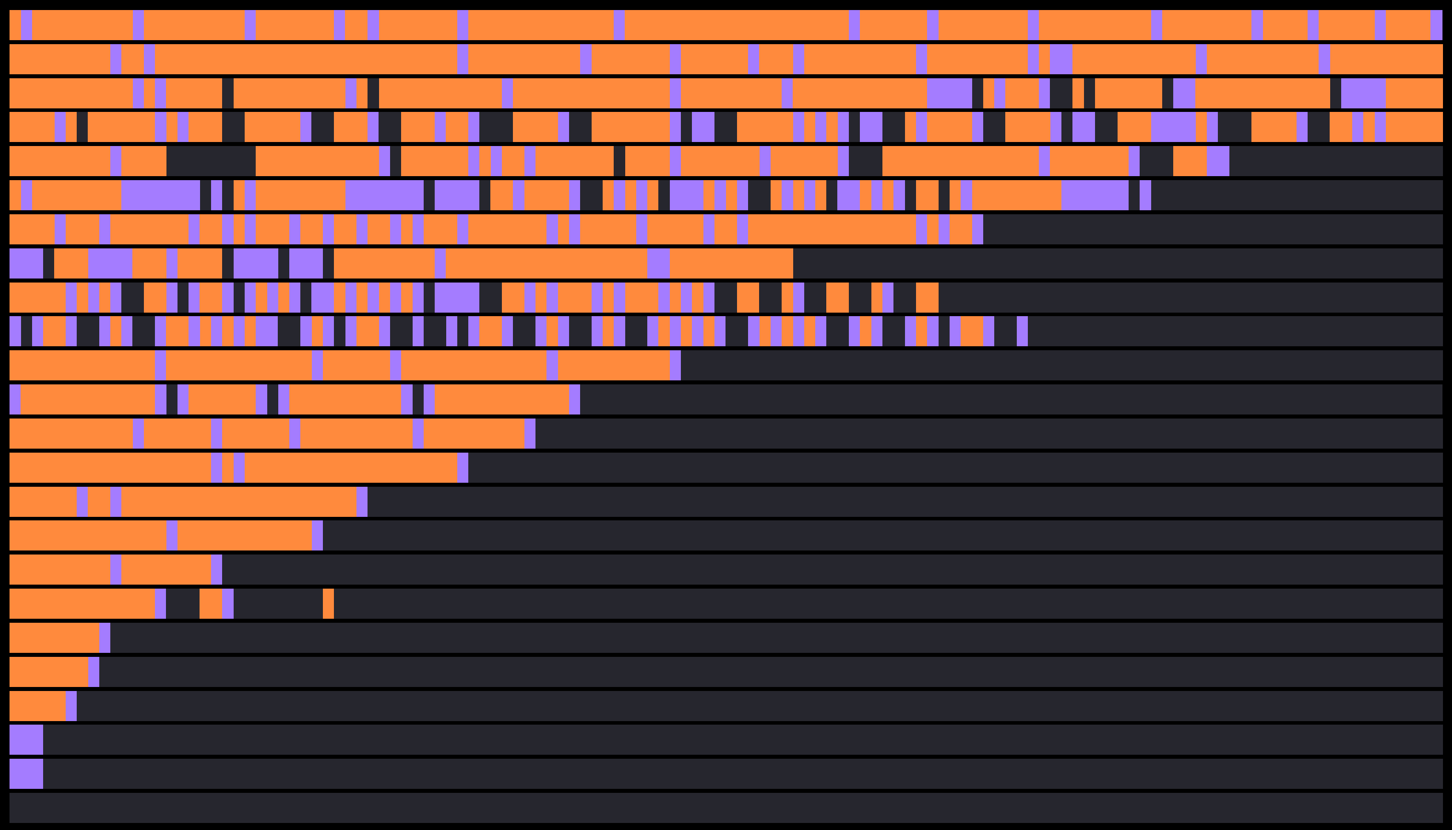

```
Okay, I need to figure out
how to handle this. The user
is talking about Outlook…
```

Letters: 41% of all tokens

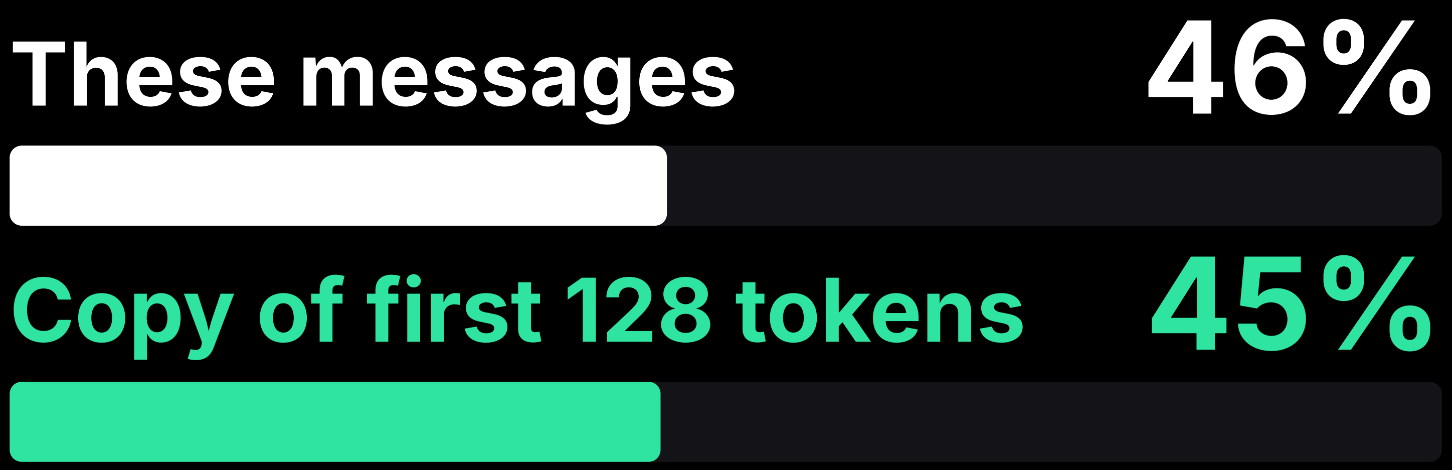

These messages 46%

Copy of first 128 tokens 45%

Strong penalty on letters, 21 more steps

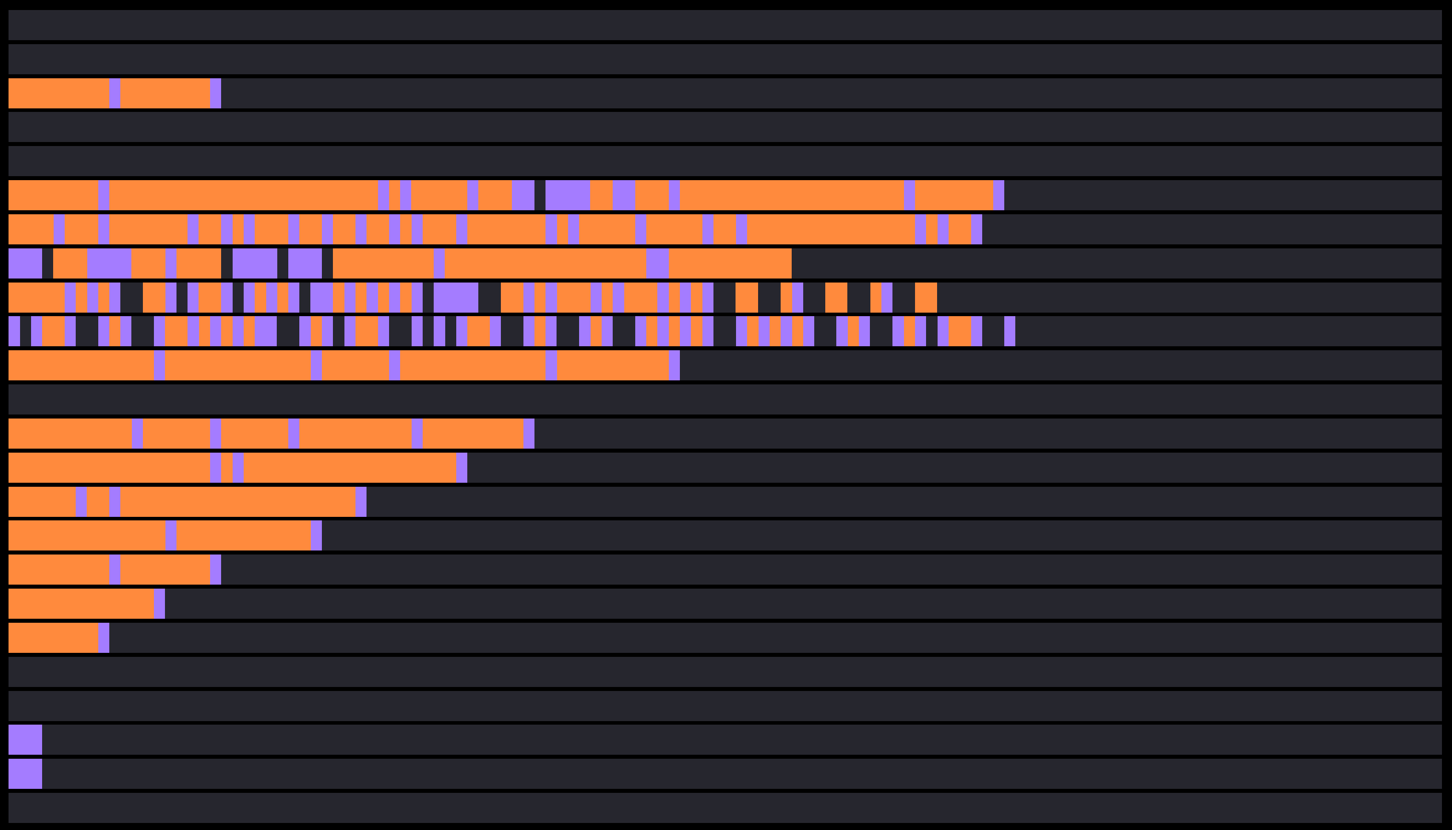

```
↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵
↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵
↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ …
```

Letters: 38% of all tokens

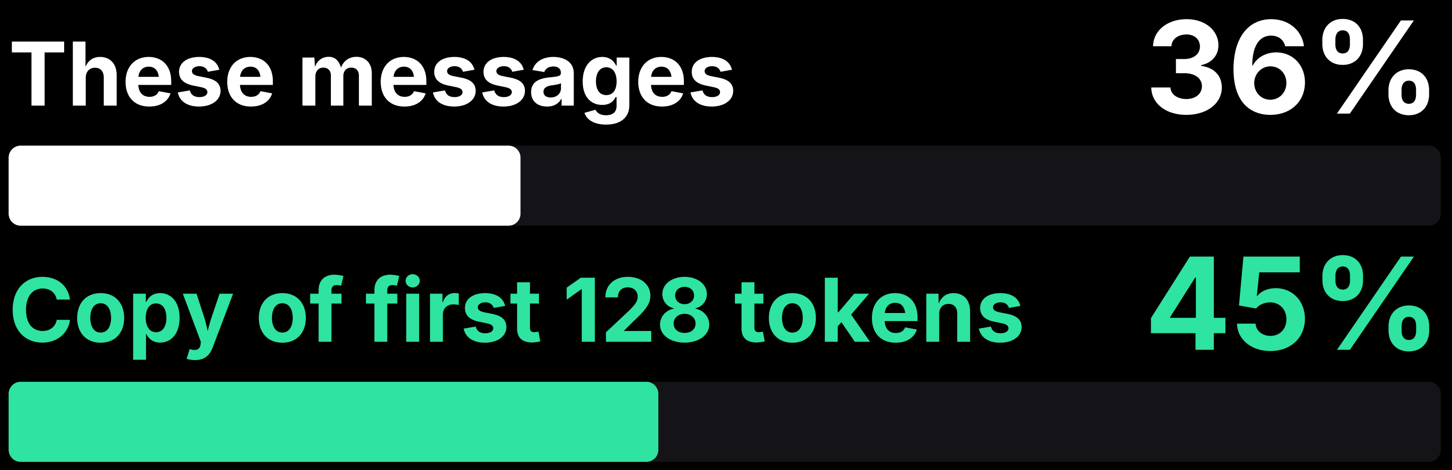

These messages 36%

Copy of first 128 tokens 45%

Strong penalty on letters and blank lines, 33 more steps

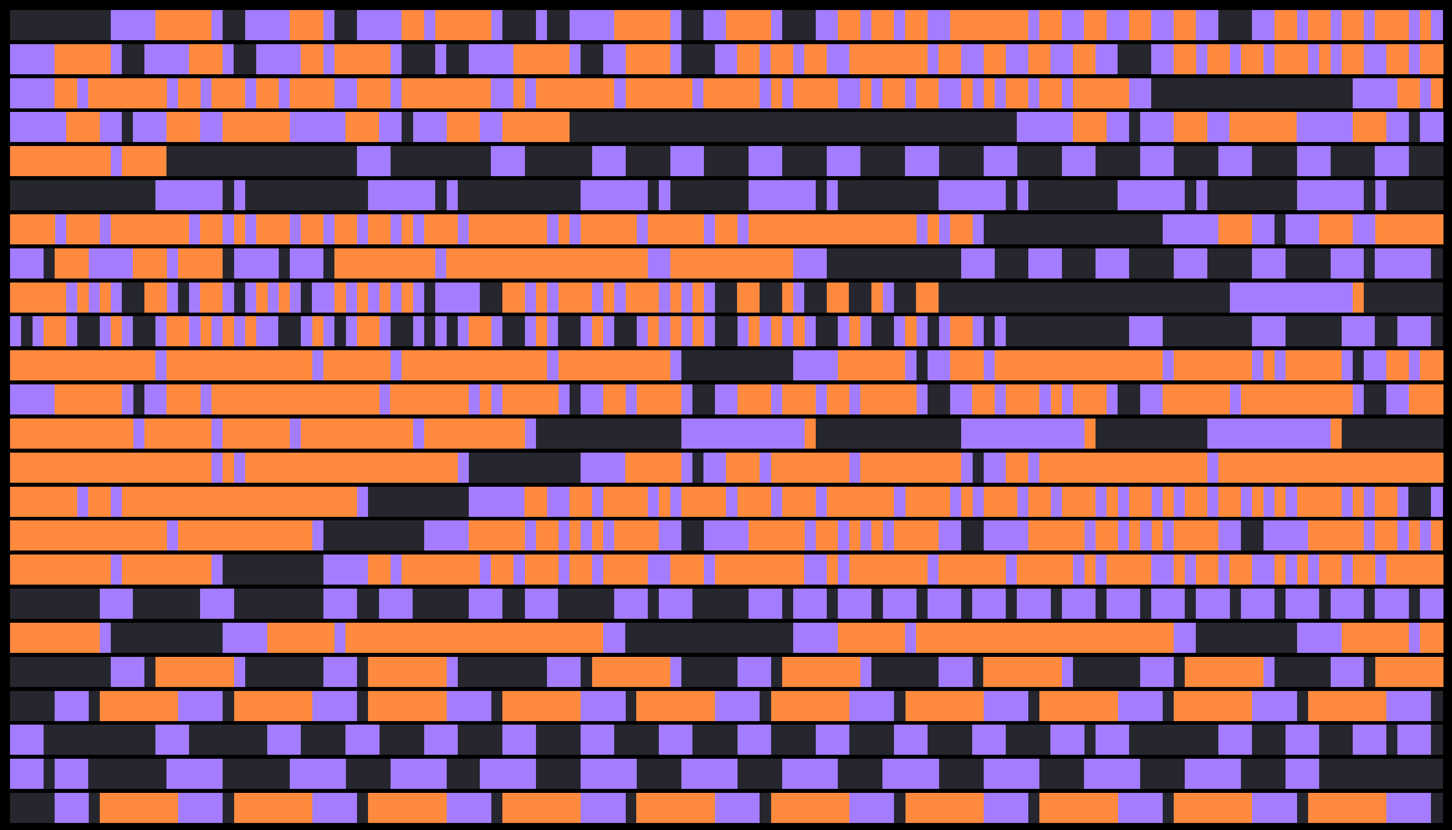

```
↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ ↵ (1) **View
Modes and Column Addition**
↵ (2) **Preview Pane…
```

Letters: 43% of all tokens

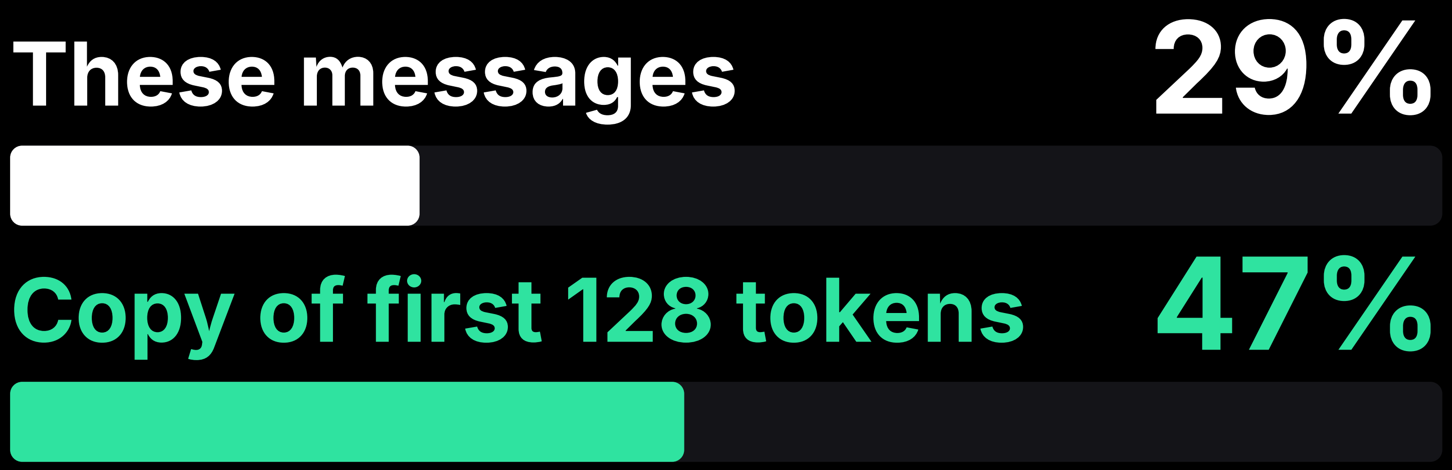

These messages 29%

Copy of first 128 tokens 47%

Bars: how close the answers get. 0%: no message. 100%: same answers as reading the whole chat.

Two copies of Qwen3-1.7B with trained adapters. Scores over 128 chats, documents and code tasks not used in training. Both strong penalties start from the weights on the left.